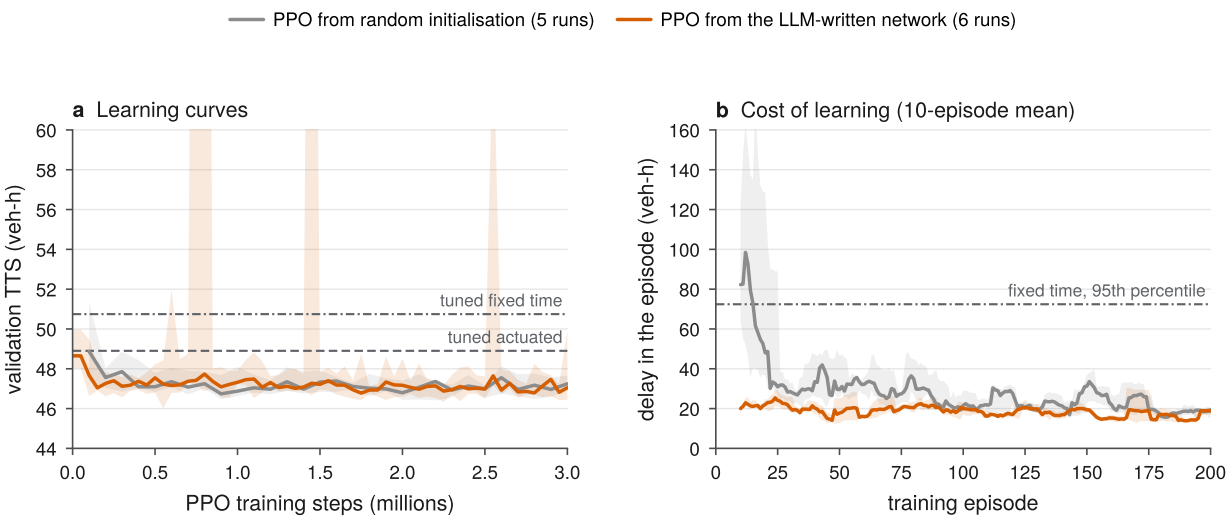
PPO from random initialisation (5 runs)
PPO from the LLM-written network (6 runs)
a Learning curves
validation TTS (veh-h)
PPO training steps (millions)
tuned fixed time
tuned actuated
b Cost of learning (10-episode mean)
delay in the episode (veh-h)
training episode
fixed time, 95th percentile